

# CSCE 5610

## Computer System Architecture

---

### The Processor

### Content

---

- Single-cycle implementation
  - All operations take the same amount of time—a single cycle.
- Multicycle implementation
  - Allows faster operations to take less time than slower ones, so overall performance can be increased.
- Pipelining
  - Lets a processor overlap the execution of several instructions, potentially leading to big performance gains.

# Single-cycle implementation

---

- We will describe the implementation a simple MIPS-based instruction set supporting just the following operations.

Arithmetic:	add	sub	and	or	slt
Data Transfer:	lw	sw			
Control:	beq				


- Today we'll build a **single-cycle implementation** of this instruction set.
  - All instructions will execute in the same amount of time; this will determine the clock cycle time for our performance equations.
  - We'll explain the datapath first, and then make the control unit.

---

## Computers are state machines

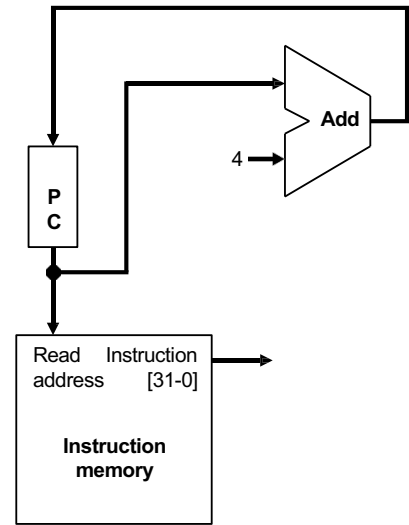
---

- A computer is just a big fancy **state machine**.
  - Registers, memory, hard disks and other storage form the state.
  - The processor keeps reading and updating the state, according to the instructions in some program.



# Instruction fetching


- The CPU is always in an infinite loop, fetching instructions from memory and executing them.
- The **program counter** or **PC** register holds the address of the current instruction.



## Basic MIPS implementation

### The first two steps for every instruction:

- Send PC to instruction memory
- Read one or two registers using fields in the machine code

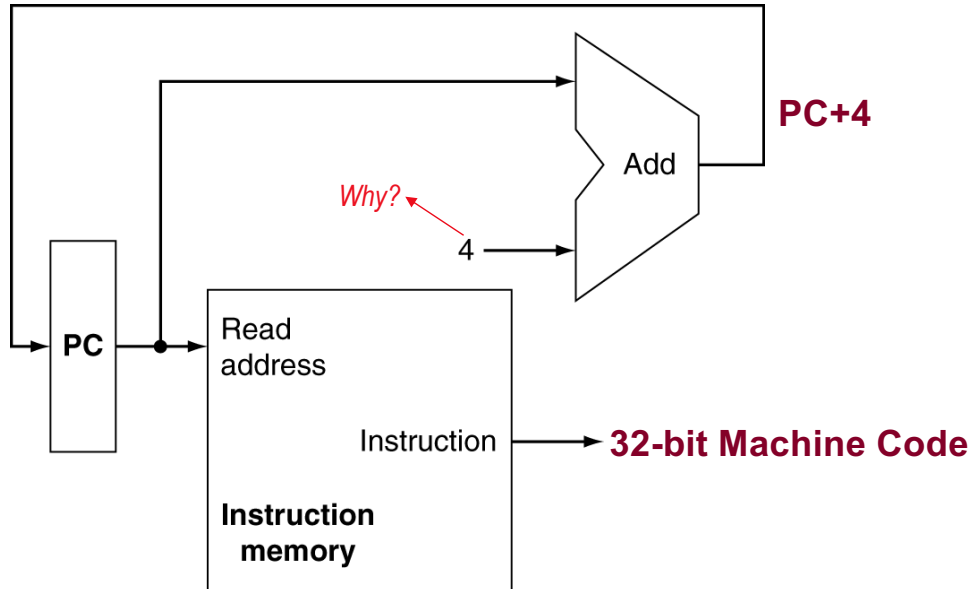


- **Instruction Memory:** stores the code and supply instruction given an address
- **Program Counter (PC):** holds the address of the current instruction
- **Adder:** increments the PC to the address of the next instruction

# Building a datapath: Fetch

## To execute any instruction:

1. **Fetch** the instruction from memory
2. Increment the PC to point at the next instruction



## Encoding R-type instructions

- Last lecture, we saw encodings of MIPS instructions as 32-bit values.
- Register-to-register arithmetic instructions use the **R-type** format.
  - **op** is the instruction opcode, and **func** specifies a particular arithmetic operation.
  - **rs**, **rt** and **rd** are source and destination registers.

op	rs	rt	rd	shamt	func
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

- An example instruction and its encoding:

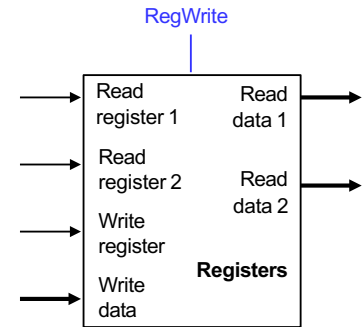
add \$20, \$9, \$10

000000	01001	01010	10100	00000	1000000
--------	-------	-------	-------	-------	---------

# Registers and ALUs


- R-type instructions must access registers and an ALU.

- Our **register file** stores thirty-two 32-bit values.
  - Each register specifier is 5 bits long.
  - You can read from two registers at a time.
  - RegWrite** is 1 if a register should be written.



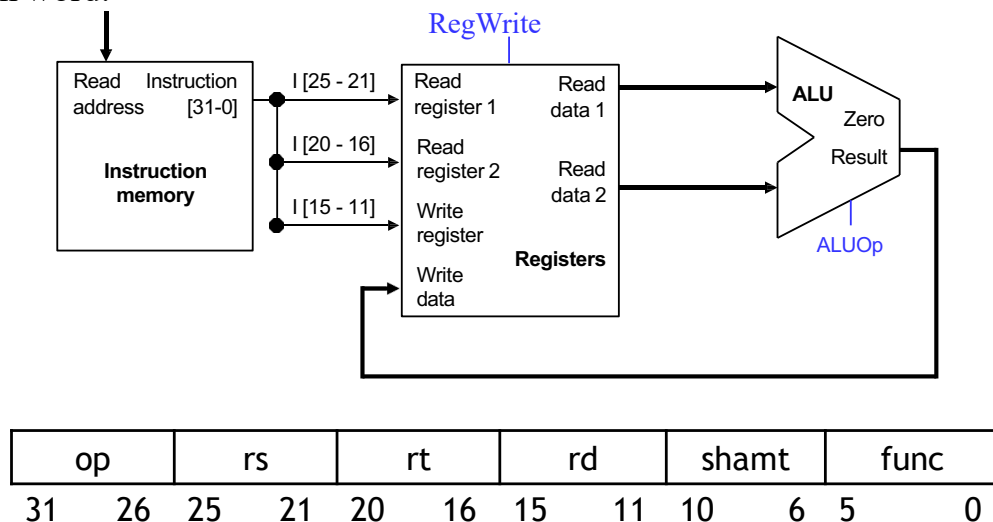
- Here's a simple **ALU** with five operations, selected by a 3-bit control signal **ALUOp**.

ALUOp	Function
000	and
001	or
010	add
110	subtract
111	slt



## Executing an R-type instruction

- Read an instruction from the **instruction memory**.
- The source registers, specified by instruction fields **rs** and **rt**, should be read from the **register file**.
- The **ALU** performs the desired operation.
- Its result is stored in the **destination register**, which is specified by field **rd** of the instruction word.



# Encoding I-type instructions

- The lw, sw and beq instructions all use the **I-type** encoding.
  - **rt** is the *destination* for lw, but a *source* for beq and sw.
  - **address** is a 16-bit **signed** constant.

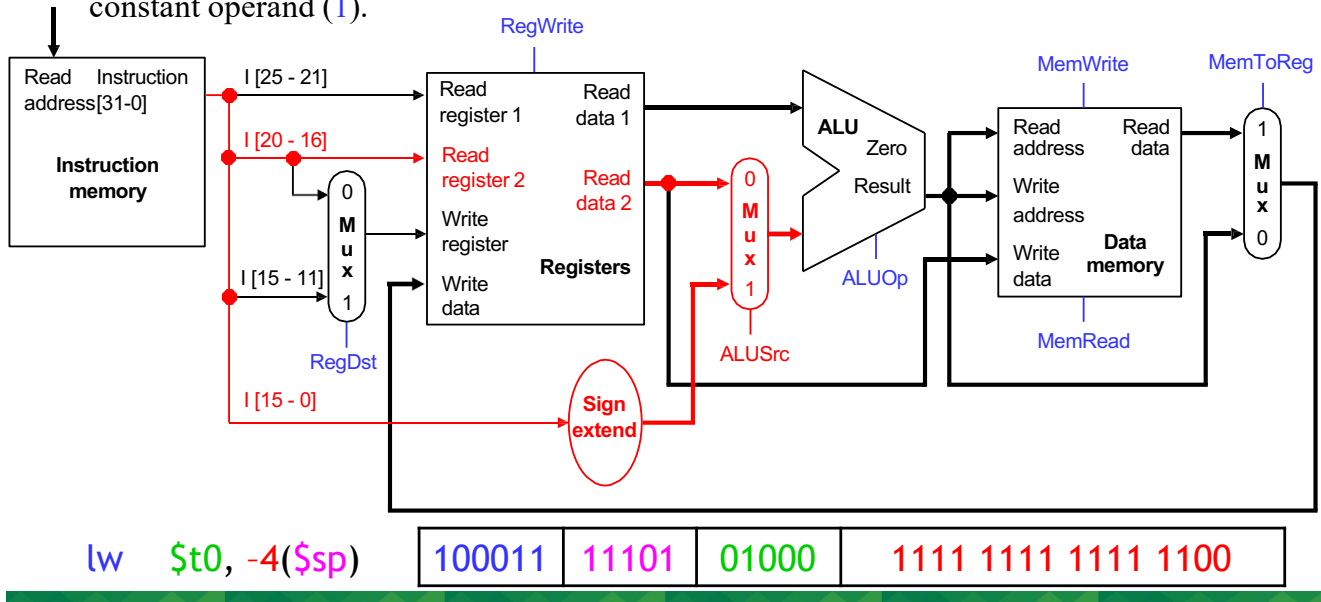
op	rs	rt	address
6 bits	5 bits	5 bits	16 bits

- Two example instructions:

lw	\$t0, -4(\$sp)	100011	11101	01000	1111 1111 1111 1100
sw	\$a0, 16(\$sp)	101011	11101	00100	0000 0000 0001 0000

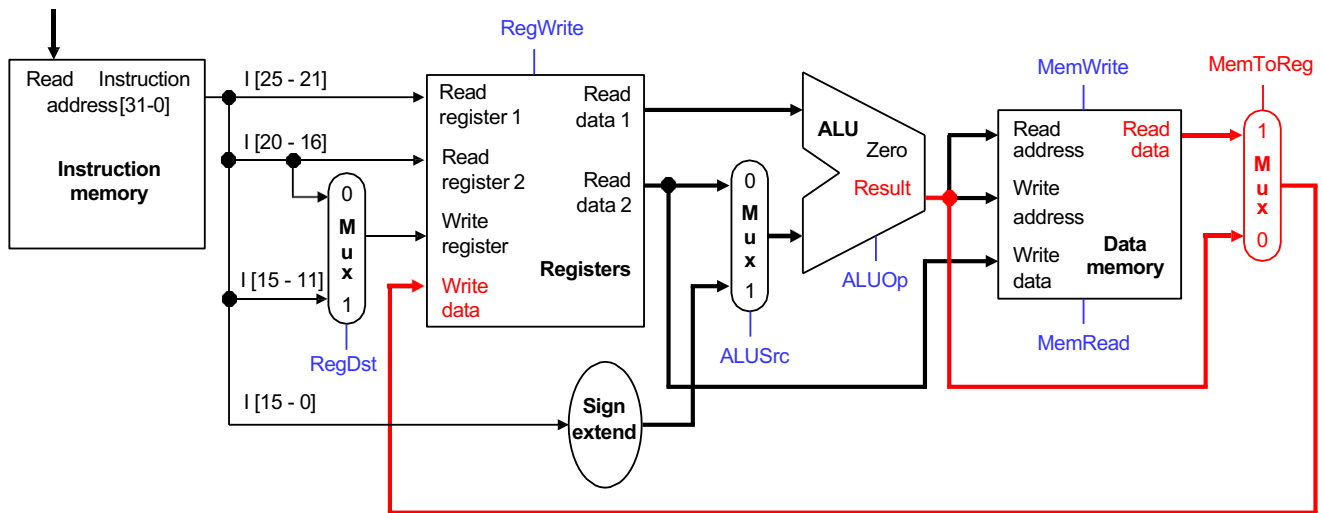
## Accessing data memory

- For an instruction like lw \$t0, -4(\$sp), the base register \$sp is added to the *sign-extended* constant to get a data memory address.
- This means the ALU must accept *either* a register operand for arithmetic instructions, *or* a sign-extended immediate operand for lw and sw.
- We'll add a multiplexer, controlled by ALUSrc, to select either a register operand (0) or a constant operand (1).




# MemToReg

- The register file's — "Write data" input has a similar problem. It must be able to store *either* the ALU output of R-type instructions, *or* the data memory output for lw.
- We add a mux, controlled by **MemToReg**, to select between saving the ALU result (0) or the data memory output (1) to the registers.

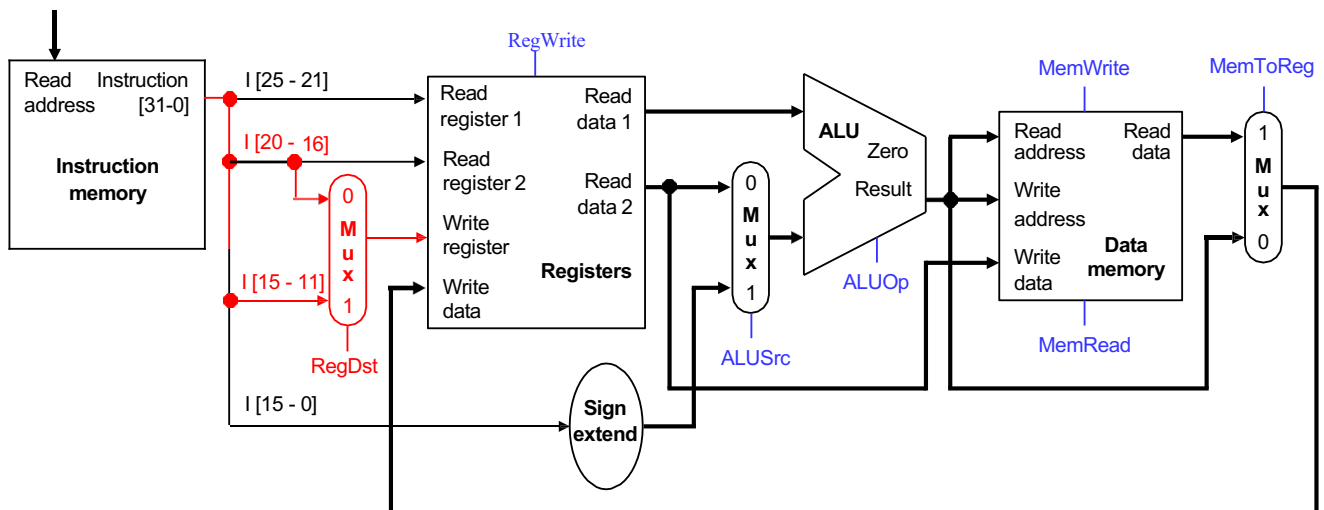


## RegDst

- A final annoyance is the destination register of lw is in *rt* instead of *rd*.



- We'll add one more mux, controlled by `RegDst`, to select the destination register from either instruction field `rt` (0) or field `rd` (1).



# Branches

- For branch instructions, the constant is not an address but an *instruction offset* from the current program counter to the desired address.

```
beq  $a1, $0, L
add  $v1, $v0, $0
add  $v1, $v1, $v1
j    Somewhere
L:   add $v1, $v0, $v0
```

- The target address L is three *instructions* past the `beq`, so the encoding of the branch instruction has 0000 0000 0000 0011 for the address field.

000100	00001	00000	0000 0000 0000 0011
op	rs	rt	address

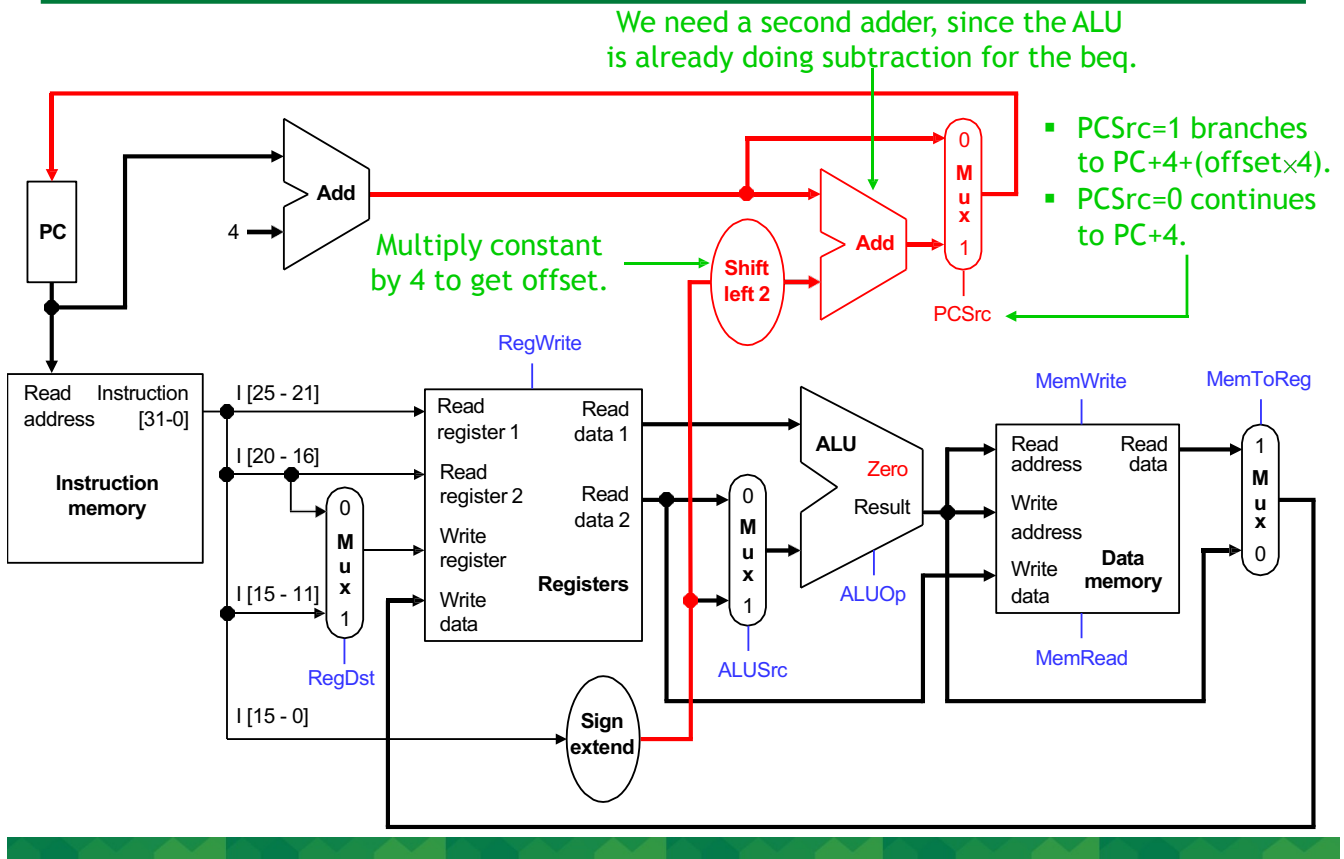
- Instructions are four bytes long, so the actual memory offset is 12 bytes.

## The steps in executing a beq

- Fetch the instruction, like `beq $a1, $0, offset`, from memory.
- Read the source registers, `$a1` and `$0`, from the register file.
- Compare the values by subtracting them in the ALU.
- If the subtraction result is 0, the source operands were equal and the PC should be loaded with the target address,  $PC + 4 + (\text{offset} \times 4)$ .
- Otherwise the branch should not be taken, and the PC should just be incremented to  $PC + 4$  to fetch the next instruction sequentially.



# Branching hardware

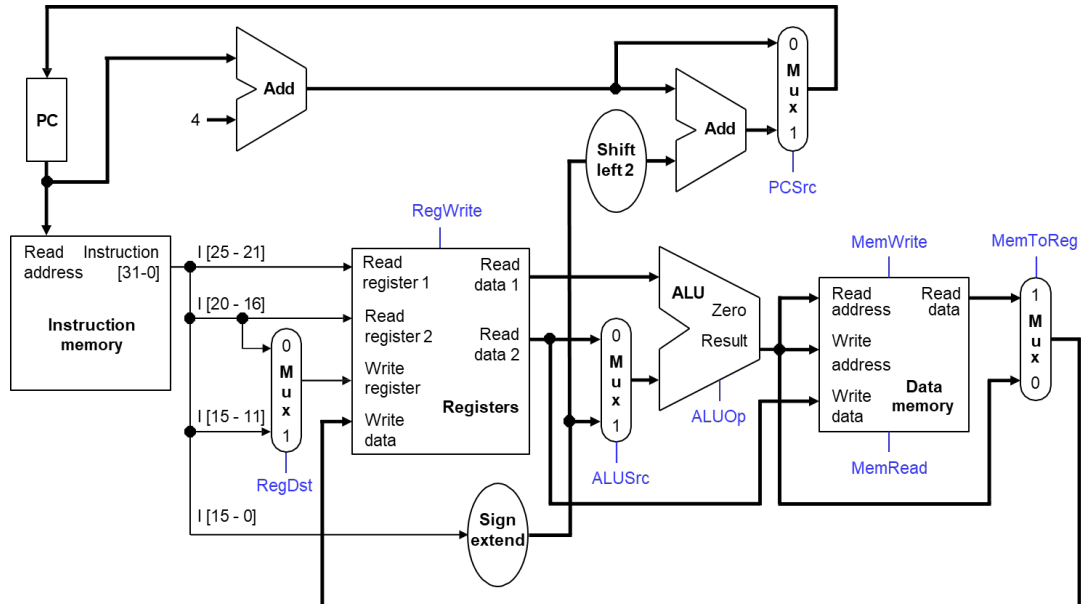


## Control

- The **control unit** is responsible for setting all the control signals so that each instruction is executed properly.
  - The control unit's input is the 32-bit instruction word.
  - The outputs are values for the blue control signals in the datapath.
- Most of the signals can be generated from the instruction **opcode** alone, and not the entire 32-bit word.
- To illustrate the relevant control signals, we will show the route that is taken through the datapath by R-type, lw, sw and beq instructions.

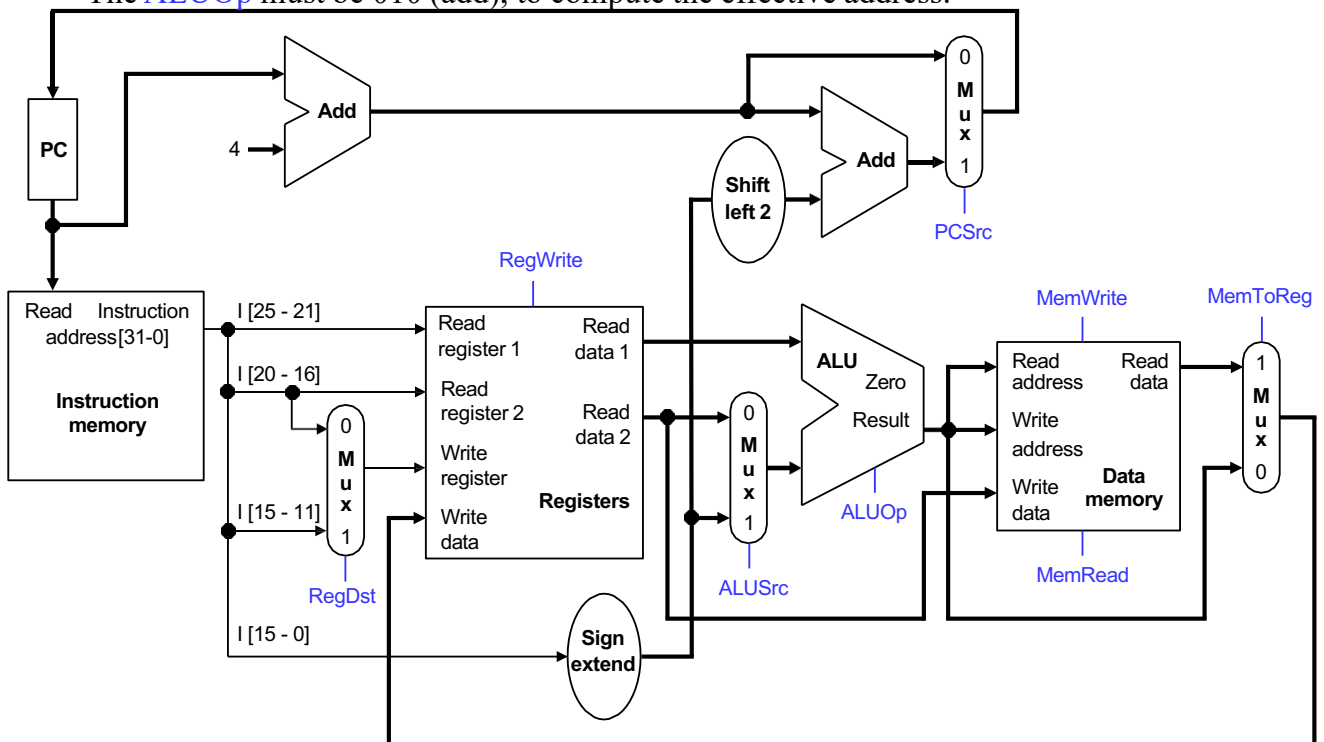
## R-type instruction path

- The R-type instructions include **add**, **sub**, **and**, **or**, and **slt**.
- The **ALUOp** is determined by the instruction's —"func" field.



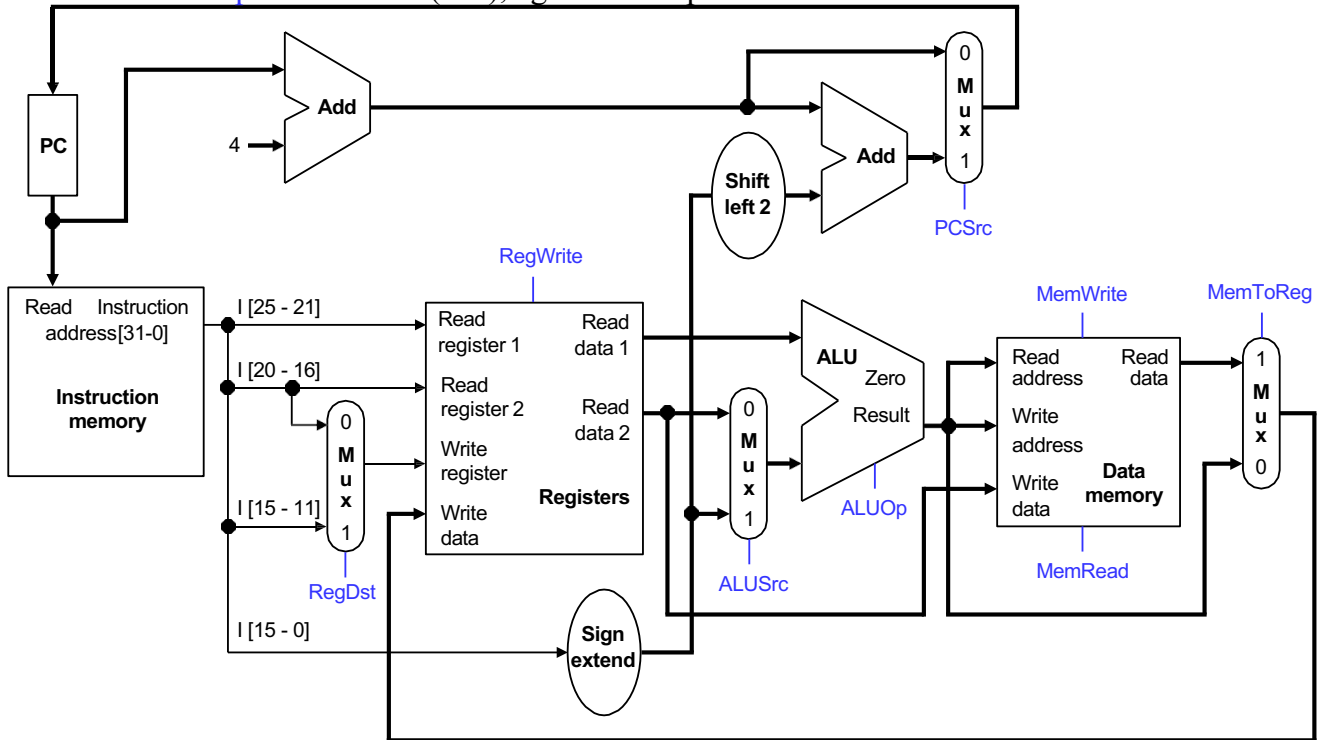
## lw instruction path

- An example load instruction is `lw $t0, -4($sp)`.
- The **ALUOp** must be 010 (add), to compute the effective address.



# sw instruction path

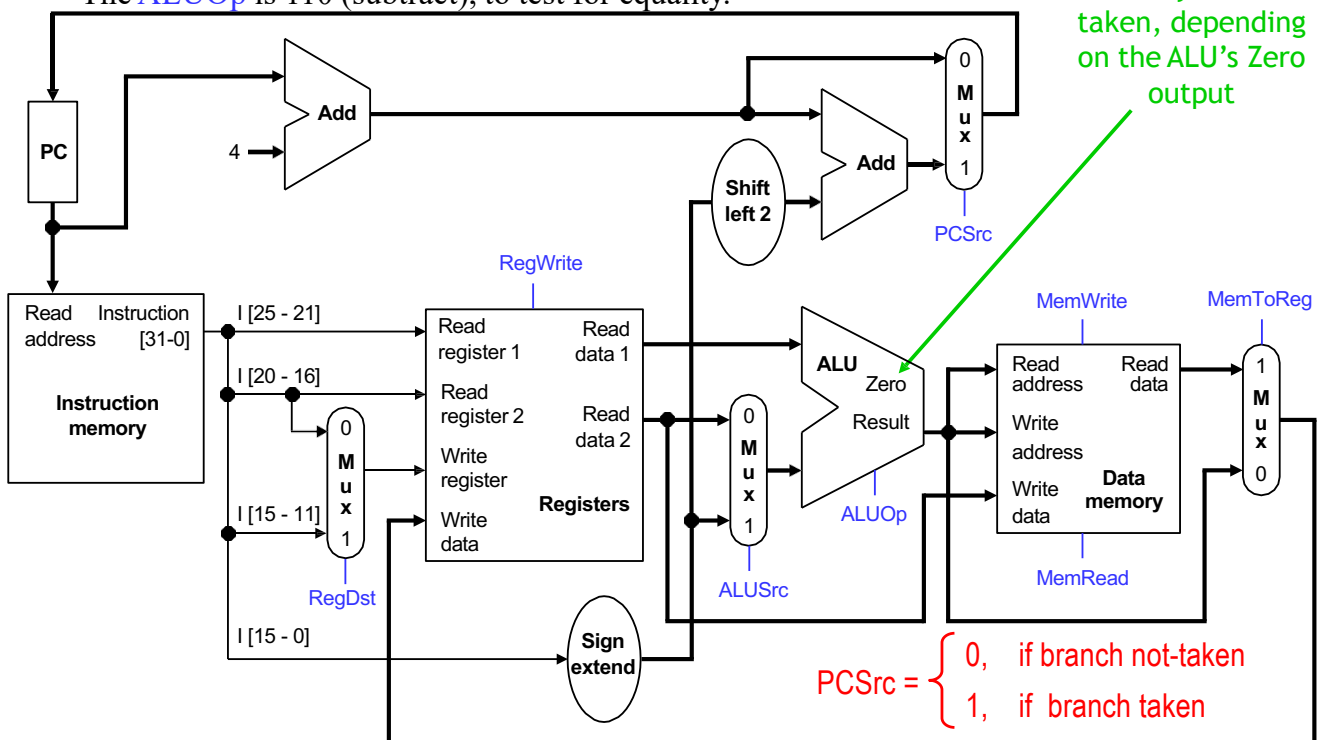
- An example store instruction is `sw $a0, 16($sp)`.
- The **ALUOp** must be 010 (add), again to compute the effective address.



# beq instruction path

- One sample branch instruction is `beq $at, $0, offset`.
- The **ALUOp** is 110 (subtract), to test for equality.

The branch may or may not be taken, depending on the ALU's Zero output



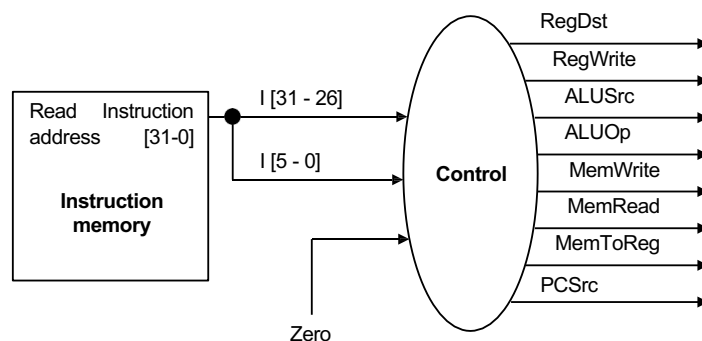
# Control Signal Table

Operation	RegDst	RegWrite	ALUSrc	ALUOp	MemWrite	MemRead	MemToReg
add	1	1	0	010	0	0	0
sub	1	1	0	110	0	0	0
and	1	1	0	000	0	0	0
or	1	1	0	001	0	0	0
slt	1	1	0	111	0	0	0
lw	0	1	1	010	0	1	1
sw	X	0	1	010	1	0	X
beq	X	0	0	110	0	0	X

- *sw* and *beq* are the only instructions that do not write any registers.
- *lw* and *sw* are the only instructions that use the constant field. They also depend on the ALU to compute the effective memory address.
- ALUOp for R-type instructions depends on the instructions' func field.
- The PCSrc control signal (not listed) should be set if the instruction is *beq* and the ALU's Zero output is true.

## Generating Control Signals

- The control unit needs 13 bits of inputs.
  - Six bits make up the instruction's opcode.
  - Six bits come from the instruction's func field.
  - It also needs the Zero output of the ALU.
- The control unit generates 10 bits of output, corresponding to the signals mentioned on the previous page.




# Single-Cycle Performance

- We saw a MIPS single-cycle datapath and control unit.
- Then, we'll explore factors that contribute to a processor's **execution time**, and specifically at the performance of the single-cycle machine.
- Next, we'll explore how to improve on the single cycle machine's performance using **pipelining**.

## Processor clock period

the processor clock oscillates between **high** and **low** signal levels:




the clock synchronizes the processor's operations on **rising** & **falling** edges:

- a register updates contents upon arrival of clock edge
- a processor **starts instruction execution** at clock edge

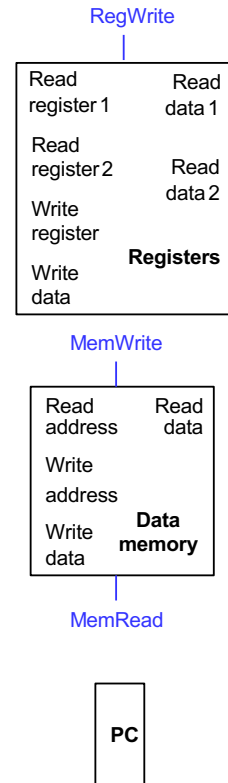
### Clock Period = Clock Cycle Time

- interval of time between 2 adjacent rising edges
- duration of 1 cycle of the clock signal
- clock period has units of seconds (nanoseconds, microseconds, or picoseconds)



# Edge-triggered State Elements


- In an instruction like `add $t1, $t1, $t2`, how do we know \$t1 is not updated until after its original value is read?
- We'll assume that our state elements are **positive edge triggered**, and are updated only on the positive edge of a clock signal.
  - The register file and data memory have explicit write control signals, **RegWrite** and **MemWrite**. These units can be written to only if the **control signal is asserted** and there is a **positive clock edge**.
  - In a single-cycle machine the PC is updated on each clock cycle, so we don't bother to give it an explicit write control signal.



## Processor Frequency

**Clock Frequency** =  $1 / (\text{Clock Period})$

a CPU's Clock Frequency is the inverse of it's Clock Cycle time



- the number of clock cycles occurring in 1 second
- measured in units of **Hertz (Hz)**
- also called the **Clock Rate**

$$F = 1 / \text{Period}$$

10 nsec clock cycle period → 100 MHz clock rate

5 nsec clock cycle period → 200 MHz clock rate

2 nsec clock cycle period → 500 MHz clock rate

1 nsec ( $10^{-9}$ ) clock cycle period → 1 GHz ( $10^9$ ) clock rate

500 psec clock cycle period → 2 GHz clock rate

250 psec clock cycle period → 4 GHz clock rate

200 psec clock cycle period → 5 GHz clock rate

When designing processors,  
we work in units of:

**Nanoseconds**

or

**GHz**

# CPU Time

---

$$CPU\ time = CPU\ Clock\ Cycles * Cycle\ Time = \frac{Clock\ Cycles}{Clock\ Rate}$$

- Performance improved by
  - Reducing number of clock cycles
  - Increasing clock rate
  - Hardware designer must often trade off clock rate against cycle count

---

## CPU Time Exercise

---

$$CPU\ time = CPU\ Clock\ Cycles * Cycle\ Time = \frac{Clock\ Cycles}{Clock\ Rate}$$

- Computer A: 2GHz clock, 10s CPU time
- Designing Computer B
  - Aim for 6s CPU time
  - Can do faster clock, but causes 1.2× clock cycles
- How fast must Computer B clock be?

$$\frac{Clock\ Rate_B}{Clock\ Rate_A} = ?$$

Let's break this down step by step:

### Computer A Information:

- **Clock rate** of Computer A: 2 GHz (which is  $2 \times 10^9$  cycles per second)
- **CPU time** for Computer A: 10 seconds

### Formula to calculate CPU time:

$$\text{CPU Time} = \text{CPU Clock Cycles} \times \text{Cycle Time}$$

Where:

- **CPU Clock Cycles** = Total number of cycles for a task to complete.
- **Cycle Time** = Time for each clock cycle, which is the reciprocal of the clock rate.

### Given for Computer A:

$$\text{CPU Time for A} = 10 \text{ seconds}$$

$$\text{Clock Rate for A} = 2 \text{ GHz} = 2 \times 10^9 \text{ cycles per second}$$

#### Step 1: Calculate total clock cycles for Computer A

We can calculate the total number of clock cycles for Computer A:

$$\text{Clock Cycles for A} = \text{Clock Rate for A} \times \text{CPU Time for A}$$

$$\text{Clock Cycles for A} = (2 \times 10^9) \times 10 = 2 \times 10^{10} \text{ clock cycles}$$

### Step 2: Design Computer B

We aim for Computer B to have a **CPU time of 6 seconds**. However, we are told that the number of clock cycles for Computer B is 1.2 times greater than that of Computer A.

$$\text{Clock Cycles for B} = 1.2 \times \text{Clock Cycles for A} = 1.2 \times (2 \times 10^{10}) = 2.4 \times 10^{10} \text{ clock cycles}$$

#### Step 3: Calculate the clock rate for Computer B

Now, we can use the formula for CPU time again to find the clock rate for Computer B. We know the CPU time is 6 seconds, and we want to find the clock rate for Computer B.

$$\text{CPU Time for B} = \frac{\text{Clock Cycles for B}}{\text{Clock Rate for B}}$$

Rearranging to solve for the clock rate:

$$\text{Clock Rate for B} = \frac{\text{Clock Cycles for B}}{\text{CPU Time for B}} = \frac{2.4 \times 10^{10}}{6} = 4 \times 10^9 \text{ Hz}$$

### Final Answer:

The clock rate for Computer B must be **4 GHz** to meet the target CPU time of 6 seconds, considering the increased number of clock cycles.



# CPI-Cycles per Instruction

---

- The **average number of clock cycles per instruction**, or **CPI**, is a function of the machine and program.
  - The CPI depends on the **actual instructions** appearing in the program— a floating-point intensive application might have a higher CPI than an integer-based program.
  - It also depends on the **CPU implementation**. For example, a Pentium can execute the same instructions as an older 80486, but faster.
- We assumed each instruction took one cycle, so we had  $CPI = 1$ .
  - The CPI can be  $>1$  due to memory stalls and slow instructions.
  - The CPI can be  $<1$  on machines that execute more than 1 instruction per cycle (superscalar).

---

## CPU Time

$$CPU\ time = Instructions\ executed * CPI * Cycle\ Time$$



# Instructions Executed

---

- Instructions executed:
  - We are not interested in the **static instruction count**, or how many lines of code are in a program.
  - Instead we care about the **dynamic instruction count**, or how many instructions are actually executed when the program runs.

```
Ostrich:  li    $4, 1000
          subi  $4, $4, 1
          bne   $4, $0, ostrich
```

---

# Instructions Executed

---

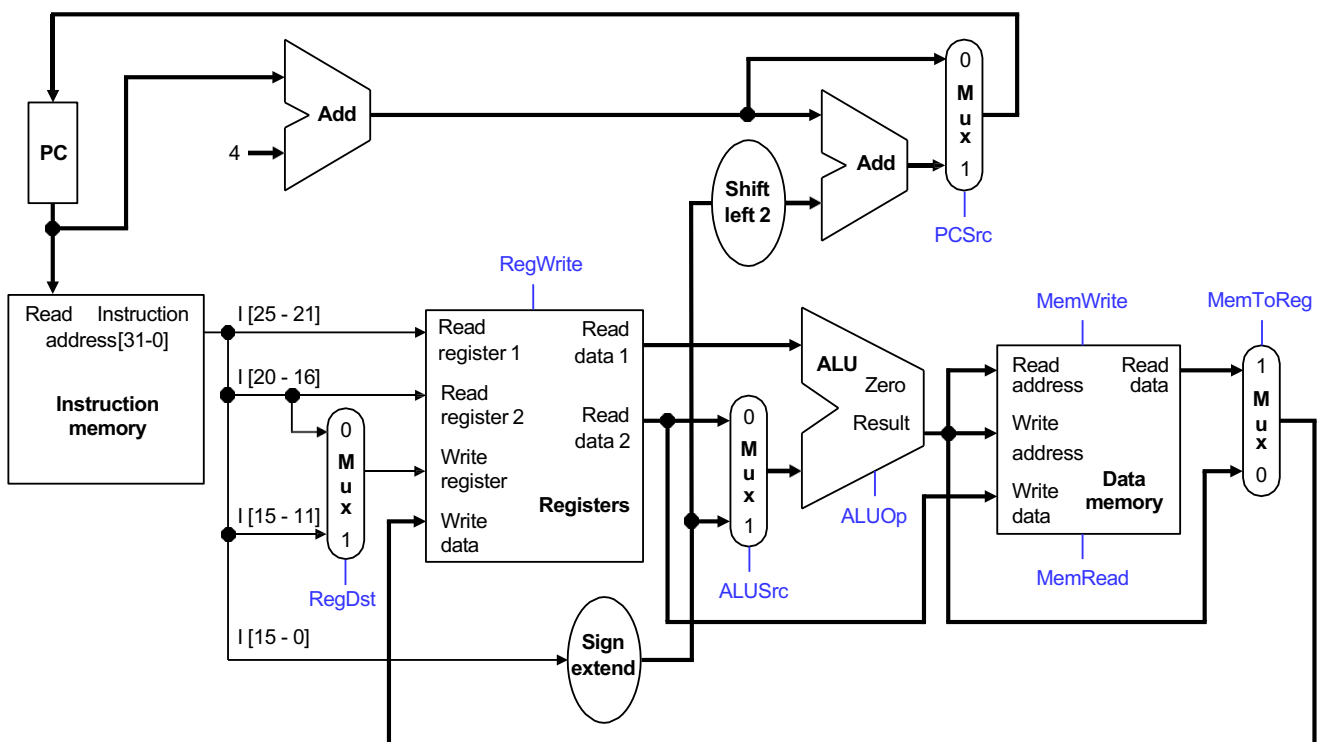
- Instructions executed:
  - We are not interested in the **static instruction count**, or how many lines of code are in a program.
  - Instead we care about the **dynamic instruction count**, or how many instructions are actually executed when the program runs.
  - There are three lines of code below, but the number of instructions executed would be 2001.

```
Ostrich:  li    $4, 1000
          subi  $4, $4, 1
          bne   $4, $0, ostrich
```

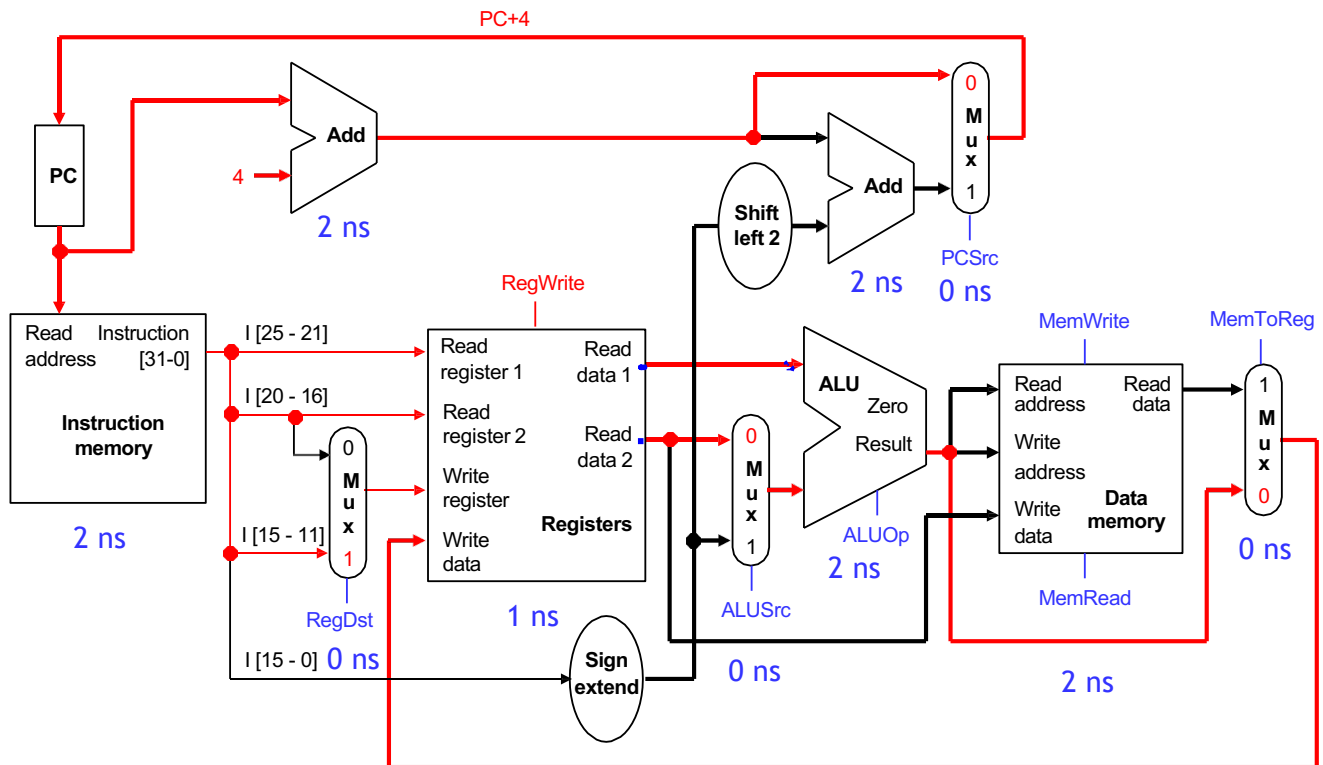
## Clock Cycle Time

- One "cycle" is the minimum time it takes the CPU to do any work.
  - The **clock cycle time** or clock period is just the length of a cycle.
  - The **clock rate**, or frequency, is the reciprocal of the cycle time.
- Generally, a higher frequency is better.
- Some examples illustrate some typical frequencies.
  - A 500MHz processor has a cycle time of 2ns.
  - A 2GHz (2000MHz) CPU has a cycle time of just 0.5ns.

## How does add go through the datapath

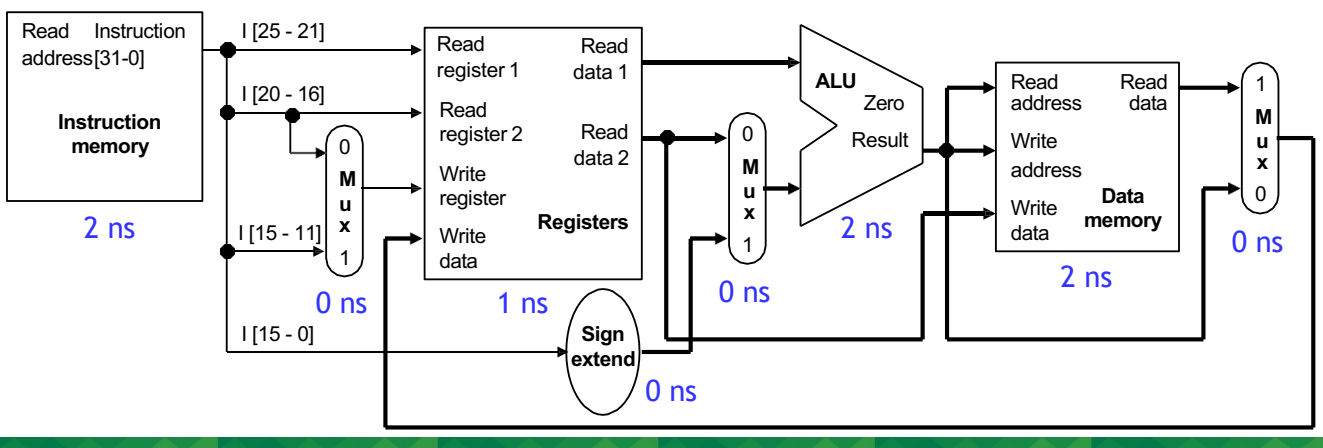


# Compute the longest path in the add instruction



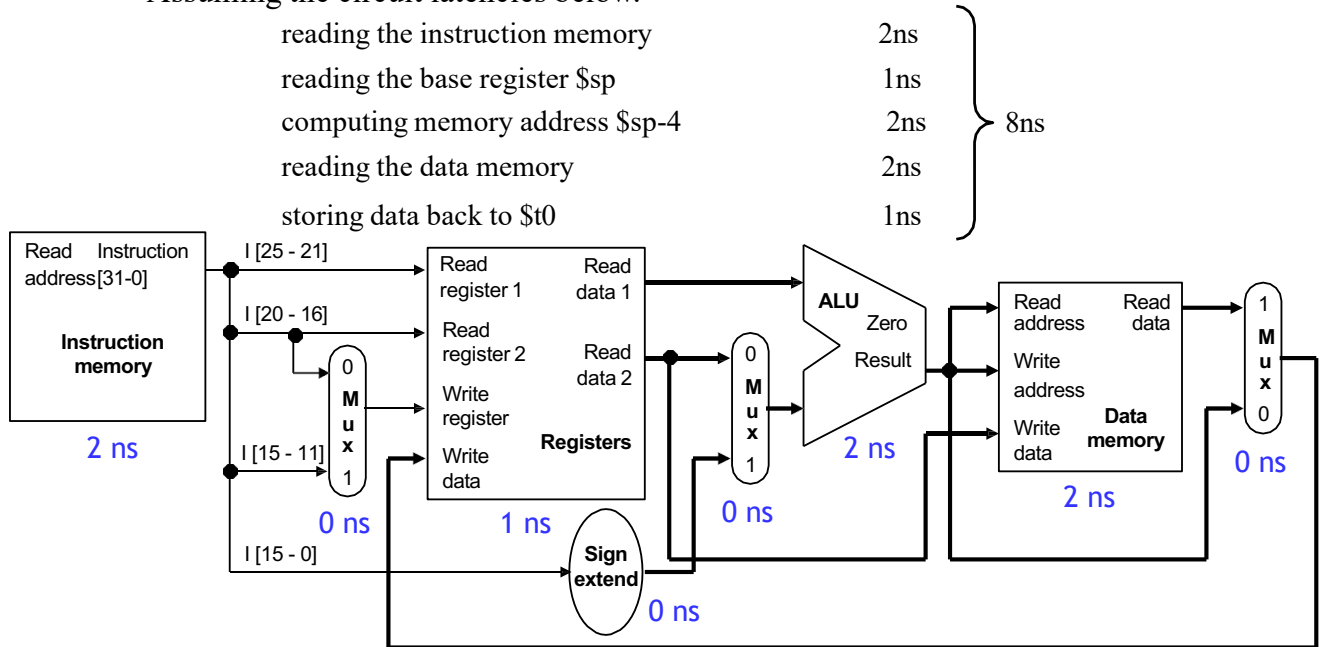
## The Slowest Instruction...

- If all instructions must complete within one clock cycle, then the cycle time has to be large enough to accommodate the *slowest* instruction.
- For example, `lw $t0, -4($sp)` is the slowest instruction needing \_\_\_\_ ns.  
— Assuming the circuit latencies below.



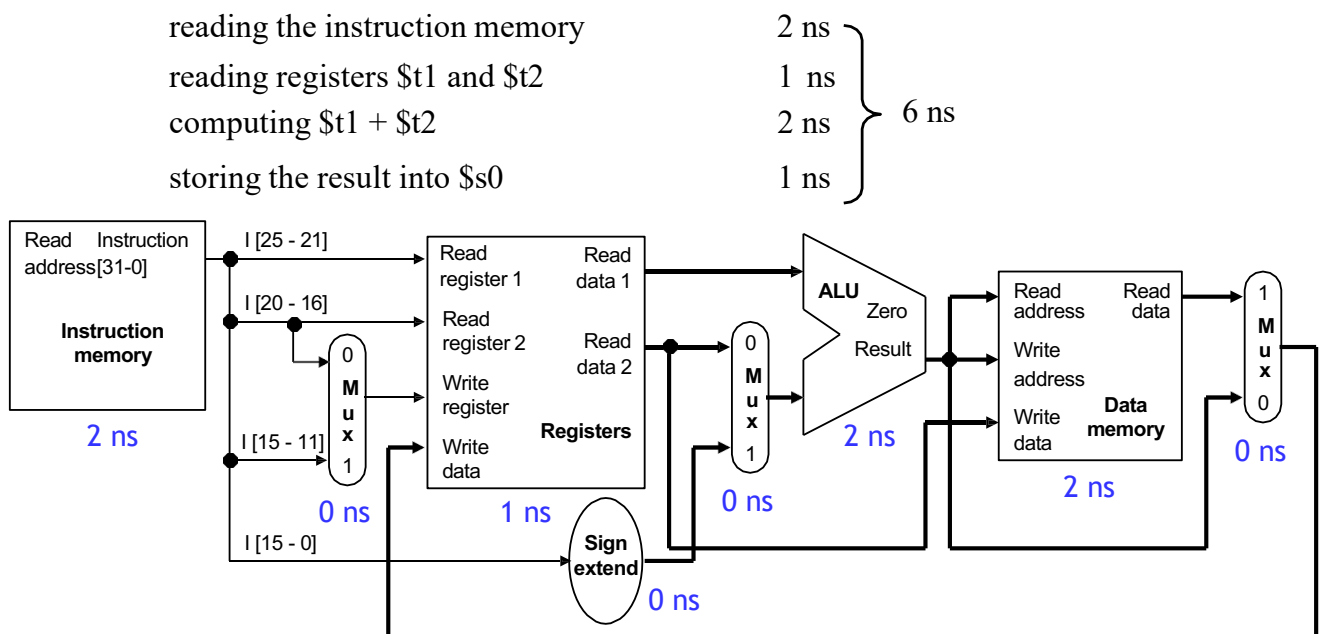
# The Slowest Instruction...

- If all instructions must complete within one clock cycle, then the cycle time has to be large enough to accommodate the *slowest* instruction.
- For example, `lw $t0, -4($sp)` is the slowest instruction needing \_\_\_\_ ns.  
— Assuming the circuit latencies below.



## ...Determines the Clock Cycle Time

- If we make the cycle time 8ns then *every* instruction will take 8ns, even if they don't need that much time.
- For example, the instruction `add $s4, $t1, $t2` really needs just 6ns.



## How bad is this?

---

- With these same component delays, a **sw** instruction would need 7ns, and **beq** would need just 5ns.
- Let's consider the **gcc** instruction mix.

Instruction	Frequency
Arithmetic	48%
Loads	22%
Stores	11%
Branches	19%



- With a single-cycle datapath, each instruction would require 8ns.
- But if we could execute instructions as fast as possible, the average time per instruction for gcc would be?

---

## How bad is this?

---

- With these same component delays, a **sw** instruction would need 7ns, and **beq** would need just 5ns.
- Let's consider the **gcc** instruction mix.

Instruction	Frequency
Arithmetic	48%
Loads	22%
Stores	11%
Branches	19%



- With a single-cycle datapath, each instruction would require 8ns.
- But if we could execute instructions as fast as possible, the average time per instruction for gcc would be?

$$(48\% \times 6\text{ns}) + (22\% \times 8\text{ns}) + (11\% \times 7\text{ns}) + (19\% \times 5\text{ns}) = 6.36\text{ns}$$

- The single-cycle datapath is about 1.26 times slower!
-

