

# Assignment-2\_Machine\_Learning\_Valapadasu\_UdayBhaskar

July 1, 2024

```
[37]: #Assignment 2: Machine Learning
      #Name: Uday Bhaskar Valapadasu
      #ID: 11696364
```

```
[38]: #Import Statements

import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
import warnings
# Suppress the FutureWarning
warnings.simplefilter(action='ignore', category=FutureWarning)
```

```
[39]: # Using the diabetes_df.csv created from assignment - 1 & Created a Pandas
      ↪dataframe from diabetes_df.csv and named it assignment2_df

assignment2_df = pd.read_csv("diabetes_df.csv")
assignment2_df
```

```
[39]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	150	33.6	
1	1	85	66	29	150	26.6	
2	8	183	64	0	150	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
..	...	...	...	...	...	...	
763	10	101	76	48	180	32.9	
764	2	122	70	27	150	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	150	30.1	
767	1	93	70	31	150	30.4	

	DiabetesPedigreeFunction	Age	Target
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0

4	2.288	33	1
...	...	...	...
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[768 rows x 9 columns]

```
[40]: # Setup the Machine Learning Model:
      #Dividing the data into features (X) array and target (y) array.
```

```
#features array
X = assignment2_df.drop(['Target'], axis=1)
#target array
y = assignment2_df['Target']
```

```
[41]: # Splitting the dataset into 80-20, 70-30, and 60-40 ratios. (Example: 80-20,
      ↪ means, 80% training data, 20% testing data, and so on.)
```

```
#Split-1 into 80-20
X_train1, X_test1, y_train1, y_test1 = train_test_split(X, y, test_size=0.2,
      ↪ random_state=42)

#Split-2 into 70-30
X_train2, X_test2, y_train2, y_test2 = train_test_split(X, y, test_size=0.3,
      ↪ random_state=42)

#Split-3 into 60-40
X_train3, X_test3, y_train3, y_test3 = train_test_split(X, y, test_size=0.4,
      ↪ random_state=42)
```

```
[42]: # For each data split, apply logistic regression machine learning model to
      ↪ build confusion matrix and accuracy estimates.
```

```
# So, importing the necessary accordingly.
```

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, accuracy_score, roc_auc_score
```

```
[43]: # Performing Logistic Regression & Building Confusion Matrix and Accuracy for
      ↪ Split-1 Training:Test(80:20) Ratio
```

```
lr_split_1 = LogisticRegression(random_state=42, max_iter=1000)
lr_split_1.fit(X_train1, y_train1)
y_pred1 = lr_split_1.predict(X_test1)
y_pred_proba1 = lr_split_1.predict_proba(X_test1)[:, 1]
cm_split_1 = confusion_matrix(y_test1, y_pred1)
```

```

accuracy_split_1 = accuracy_score(y_test1, y_pred1)
auc_split_1 = roc_auc_score(y_test1, y_pred_proba1)

# Displaying the Confusion Matrix, Accuracy, and AUC
print("Confusion Matrix and Accuracy for Split-1 Training:Test(80:20)")
print("Confusion Matrix:")
print(cm_split_1)
print("Accuracy:", accuracy_split_1)
print("AUC:", auc_split_1)

```

Confusion Matrix and Accuracy for Split-1 Training:Test(80:20)  
 Confusion Matrix:  
 [[80 19]  
 [19 36]]  
 Accuracy: 0.7532467532467533  
 AUC: 0.8165289256198347

```

[44]: # Performing Logistic Regression & Building Confusion Matrix, Accuracy, and AUC
      ↪ for Split-2 Training:Test(70:30) Ratio
lr_split_2 = LogisticRegression(solver='lbfgs', random_state=42, max_iter=1000)
lr_split_2.fit(X_train2, y_train2)
y_pred2 = lr_split_2.predict(X_test2)
y_pred_proba2 = lr_split_2.predict_proba(X_test2)[:, 1]
cm_split_2 = confusion_matrix(y_test2, y_pred2)
accuracy_split_2 = accuracy_score(y_test2, y_pred2)
auc_split_2 = roc_auc_score(y_test2, y_pred_proba2)

# Displaying the Confusion Matrix, Accuracy, and AUC
print("Confusion Matrix, Accuracy, and AUC for Split-2 Training:Test(70:30)")
print("Confusion Matrix:")
print(cm_split_2)
print("Accuracy:", accuracy_split_2)
print("AUC:", auc_split_2)

```

Confusion Matrix, Accuracy, and AUC for Split-2 Training:Test(70:30)  
 Confusion Matrix:  
 [[121 30]  
 [30 50]]  
 Accuracy: 0.7402597402597403  
 AUC: 0.7964403973509934

```

[45]: # Performing Logistic Regression & Building Confusion Matrix, Accuracy, and AUC
      ↪ for Split-3 Training:Test(60:40) Ratio
lr_split_3 = LogisticRegression(random_state=42, max_iter=1000)
lr_split_3.fit(X_train3, y_train3)
y_pred3 = lr_split_3.predict(X_test3)
y_pred_proba3 = lr_split_3.predict_proba(X_test3)[:, 1]
cm_split_3 = confusion_matrix(y_test3, y_pred3)

```

```

accuracy_split_3 = accuracy_score(y_test3, y_pred3)
auc_split_3 = roc_auc_score(y_test3, y_pred_proba3)

# Displaying the Confusion Matrix, Accuracy, and AUC
print("Confusion Matrix, Accuracy, and AUC for Split-3 Training:Test(60:40)
↳Ratio")
print("Confusion Matrix:")
print(cm_split_3)
print("Accuracy:", accuracy_split_3)
print("AUC:", auc_split_3)

```

Confusion Matrix, Accuracy, and AUC for Split-3 Training:Test(60:40) Ratio

Confusion Matrix:

```
[[166  40]
 [ 36  66]]
```

Accuracy: 0.7532467532467533

AUC: 0.8213401865600609

[46]: *#Which data split is providing you the best accuracy?*

*#Ans: As, we observe split\_1(80-20), split\_3(60-40) ratio are on a tie giving*  
*↳the best accuracy of 75% when compared to data split\_2(70:30)*  
*# Here, I am selecting split\_1(80-20) for the further analysis.*

[47]: *# Importing all the necessary libraries*

```

from sklearn.utils import resample
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import roc_curve, roc_auc_score

```

[48]: *# Compute ROC curve and ROC area for Split-1*

```

fpr_split1, tpr_split1, thresholds_split1 = roc_curve(y_test1, y_pred_proba1)
roc_auc_split1 = auc(fpr_split1, tpr_split1)

# Bootstrap analysis for Split-1 (80:20)
n_iterations = 1000
bootstrap_accuracies = []

# Perform bootstrap sampling
for i in range(n_iterations):
    X_resampled, y_resampled = resample(X_train1, y_train1, replace=True,
↳random_state=i)
    lr_split_1.fit(X_resampled, y_resampled) # Fit logistic regression on
↳resampled data
    y_pred_resampled = lr_split_1.predict(X_test1) # Predict on test set
    accuracy_resampled = accuracy_score(y_test1, y_pred_resampled) # Calculate
↳accuracy

```

```

        bootstrap_accuracies.append(accuracy_resampled) # Store accuracy from each
        ↪ iteration

# Calculate p-value and confidence intervals
p_value_split1 = np.mean(np.array(bootstrap_accuracies) >= accuracy_split_1)
confidence_interval_split1 = np.percentile(bootstrap_accuracies, [2.5, 97.5])

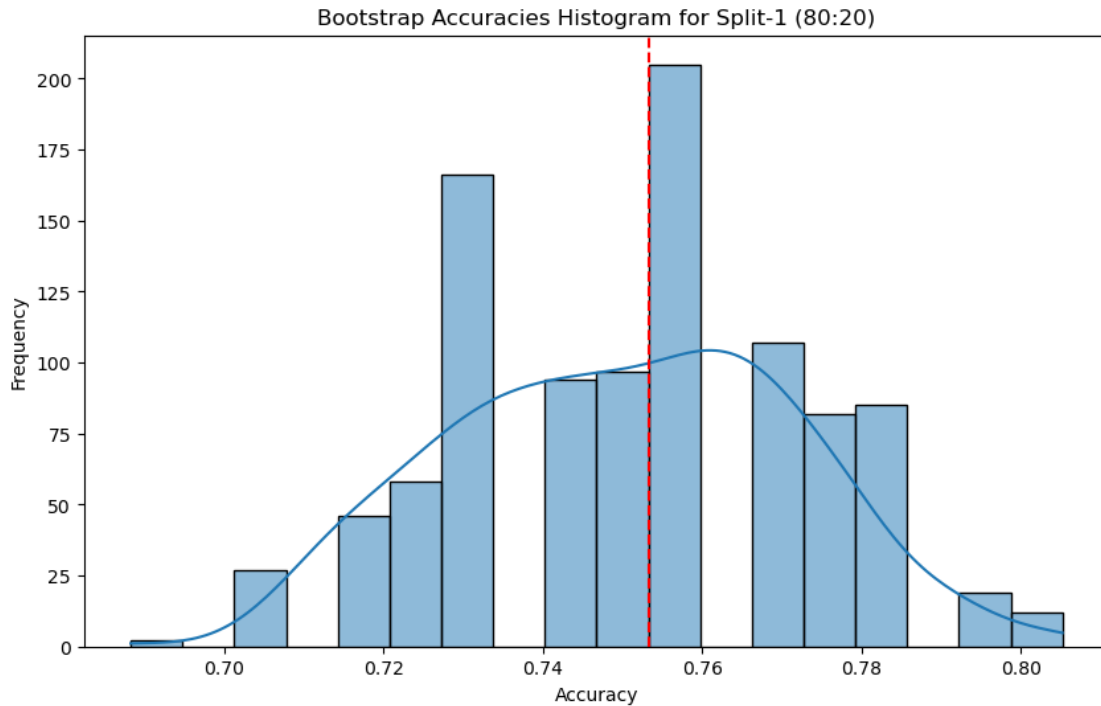
# Plot histogram of bootstrap accuracies
plt.figure(figsize=(10, 6))
sns.histplot(bootstrap_accuracies, kde=True)
plt.axvline(accuracy_split_1, color='r', linestyle='--')
plt.xlabel('Accuracy')
plt.ylabel('Frequency')
plt.title('Bootstrap Accuracies Histogram for Split-1 (80:20)')
plt.show()

# Write the report with metrics for Split-1 (80:20)
report_split1 = f"""
Selected Data Split Ratio (80:20) Metrics:
    Confusion Matrix:
{cm_split_1}
    Accuracy: {accuracy_split_1}
    AUC: {auc_split_1}
    ROC AUC: {roc_auc_split1}
    ROC Thresholds: {thresholds_split1}

Bootstrap Analysis for Split-1:
    P-Value: {p_value_split1}
    Confidence Interval: {confidence_interval_split1}
"""

print(report_split1)

```



Selected Data Split Ratio (80:20) Metrics:

Confusion Matrix:

[[80 19]

[19 36]]

Accuracy: 0.7532467532467533

AUC: 0.8165289256198347

RDC AUC: 0.8165289256198347

RDC Thresholds: [1.97062516 0.97062516 0.96040546 0.87284274 0.87075569  
0.78031972

0.77299541 0.77209963 0.76866991 0.72395837 0.70889906 0.70883013  
0.69285014 0.66792411 0.66001388 0.65390286 0.65205796 0.63552015  
0.62281485 0.61803155 0.57749256 0.56826283 0.55545133 0.54548516  
0.5319286 0.5140412 0.49408098 0.49150511 0.4436292 0.42606807  
0.40481265 0.37964985 0.36247904 0.327527 0.28621408 0.27489661  
0.26885186 0.26091845 0.26086909 0.25787497 0.23546844 0.23035508  
0.21334583 0.20356076 0.15930276 0.15109309 0.1416267 0.14056454  
0.13153163 0.12834498 0.12069508 0.11672998 0.04015114 0.0384741  
0.00415931]

Bootstrap Analysis for Split-1:

P-Value: 0.51

Confidence Interval: [0.70779221 0.79220779]